



Storage Management



Unit objectives

After completing this unit, you should be able to:

- Review of modern internal hard disk type
- Device naming convention
- Perform Disk partitioning
- Understand what and why need File System
- Create and Mount File System
- Unmount File System

Overview of modern internal hard disk type - 1/2

- Most common device for persistent storage
- Three common types: SATA, SAS, and NVMe
- Integrated Drive Electronics (IDE) - Legacy
 - Some cabling limitations
 - Device naming `/dev/hda, hdb, ..., hdh`
- Serial Advanced Technology Attachment (SATA)
 - Offers some advantages over ATA/EIDE:
 - Hot swappable, faster transfer speed, smaller cabling
- Small Computer System Interface (SCSI) - Legacy
 - Different subtypes: `fast, wide, fast wide, ultra-wide, ...`
 - Max 7 or 15 disks on one bus (depends on subtype)
 - Needs correct termination at both ends of bus
 - Generally more expensive
 - Device naming `/dev/sda, ..., sdz, sda1, ..., sddx`

Overview of modern internal hard disk type – 2/2

- Serial Attached SCSI (SAS)
 - offer better disk speed (10,000 / 15,000 rpm)
 - interface speed of upto 22.5 GB/s
 - Max upto 65,535 disks
- Non-Volatile Memory express (NVMe)
 - interface speed of upto 32 GB/s
 - offers high IOPS, low latency (≤ 20 microseconds)
 - Two variants : M.2 and U.2



Devices naming convention

- All character devices have a representation in `/dev`

```
# ls -l /dev/{sd,nvme}*
```

```
crw-----. 1 root root 243, 0 Sep 27 07:32 /dev/nvme0
```

```
brw-rw----. 1 root disk 259, 0 Sep 27 07:32 /dev/nvme0n1
```

```
brw-rw----. 1 root disk 259, 1 Sep 27 07:32 /dev/nvme0n1p1
```

```
brw-rw----. 1 root disk 259, 2 Sep 27 13:35 /dev/nvme0n1p2
```

```
brw-rw----. 1 root disk  8, 0 Sep 27 13:35 /dev/sda
```

```
brw-rw----. 1 root disk  8, 16 Sep 27 13:35 /dev/sdb
```

```
brw-rw----. 1 root disk  8, 32 Sep 27 13:34 /dev/sdc
```

```
brw-rw----. 1 root disk  8, 48 Sep 27 07:32 /dev/sdd
```

```
brw-rw----. 1 root disk  8, 64 Sep 27 07:32 /dev/sde
```

Disk partitioning principles : MBR / DOS

- Four primary partitions, one of which might be an "extended" partition
 - /dev/hda1 – /dev/hda4
- Extended partition: Large number of logical partitions
 - /dev/hda5 and up
- Limitation
 - to 2TB
 - 15 logical partitions



Disk partitioning principles → GPT

- GUID Partition Table
- UEFI firmware standard
- Address limitation in MBR-based scheme
- Encouraged to use!
- Cool stuff
 - max of 128 partitions
 - ZiB or 8 billion TB
 - GUID to identify each disk & partition
 - Redundancy of partition information



Partitioning tools

- **fdisk**
 - Virtually every PC OS comes with a tool fdisk to create partitions for that OS.
 - Linux, Windows, and so forth
- **parted**
 - GPLed Linux program, available at www.gnu.org
 - Can create/resize/move/delete partitions
- **GParted, QTParted**
 - GUI utilities that use GNU Parted
 - Can create/resize/move/delete partitions
- **Disk Druid and others**
 - Partitioning program integrated in Linux install program
- **Cockpit web-based interface**

The fdisk tool

View disk configuration

```
# fdisk -l /dev/sda
```

Configure disk's partition

```
# fdisk /dev/sda
```

Display current partition

Command (m for help): p

Create new partition

Command (m for help): n

Create delete partition

Command (m for help): d

The parted tool - interactively

View disk configuration

```
# parted /dev/sda print
```

Configure disk's partition

```
# parted /dev/sda
```

Decide on partition type

```
(parted) mklabel { gpt | msdos }
```

Display current partition

```
(parted) print
```

Create new partition

```
(parted) mkpart
```

Create delete partition

```
(parted) rm
```

Rescue lost partition near start and end

```
(parted) rescue
```

Resize partition

```
(parted) resizepart
```

The parted tool – non-interactively

- Convert disk into GPT

```
# parted /dev/sda mklabel gpt
```

- View disk configuration

```
# parted /dev/sda print
```

- Create first partition

```
# parted /dev/sda mkpart 1 xfs 1 100m
```

- Create next partition

```
# parted /dev/sda mkpart 2 ext4 101m 500m
```

- Create third partition

```
# parted /dev/sda mkpart 3 xfs 501 1g
```

The parted tool – non-interactively

- Delete partition number 2

```
# parted /dev/sda rm 2
```

- Resize partition 3 to 1GB

```
# parted /dev/sda resizepart 3 1g
```

- Make sure kernel detect new partitions and create necessary devices in /dev directory

```
# udevadm settle
```

The parted tool – manage MBR partitions

- Convert disk into MBR

```
# parted /dev/sda mklabel msdos  
destroy data
```



Warning: This will

- Resize partition 3 to 1GB

```
# parted /dev/sda resizepart 3 1g
```

Shortcuts to partitioning

- Copy partitions

```
# sfdisk -d /dev/sda | sfdisk /dev/sdb
```

- Quickly delete all partitions in /dev/sda

```
# dd if=/dev/zero of=/dev/sda bs=512 count=1
```

Verify all configured partitions

- Using fdisk

```
# fdisk -l /dev/sda
```

- Using parted

```
# parted /dev/sda print
```

- Using lsblk


```
# lsblk
```

- Using blkid

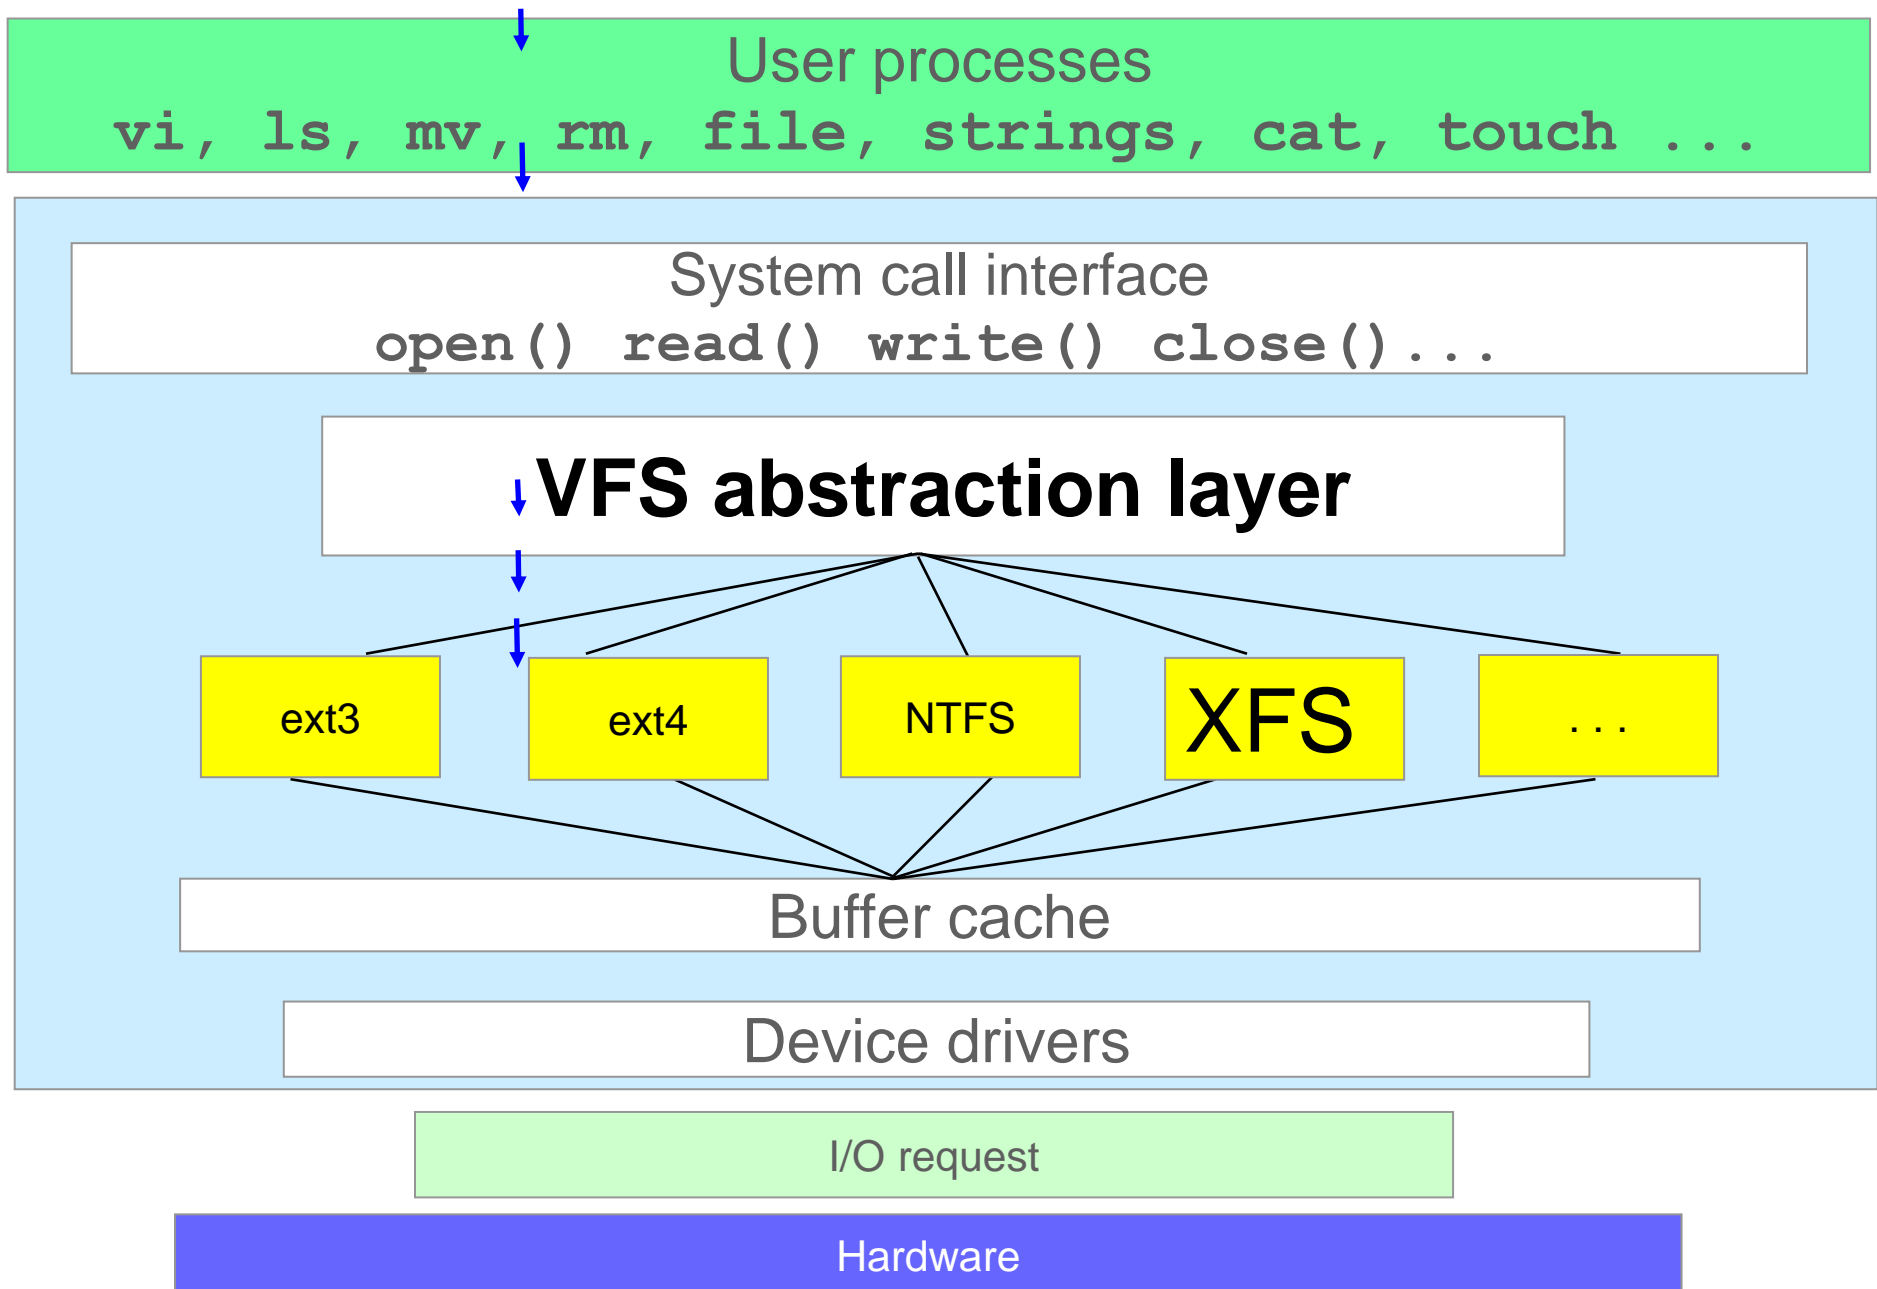
```
# blkid
```

Why need file system?

- Place to store files and refer to them
- Hierarchical structure through use of directories
- A file system can be stored on any block device
 - Floppy disk
 - Hard disk
 - Partition
 - RAID, LVM volume
 - File (for use with a loop device)
 - RAM disk



The virtual file system



Format partition with File System

- Traditional: ext2
- Second generation: ext3, ReiserFS, IBM JFS, xfs
- Next generation: ext4, GFS2, Reiser4
- FAT-12, FAT-16, FAT-32, VFAT, NTFS (read-only)
- CD-ROM (ISO 9660)
- NFS (Network File System)
- SMBFS (Windows share), NCPFS (Novell NetWare share)
- `/proc` (for kernel and process information)
- SHMFS (shared memory file system)
- GPFS, Lustre (clustering file systems)
- XFS (high performing 64-bit journaling FS)

Creating file system

- Creating a file system is done with a `mkfs` variant.
 - `mkfs.ext2`, `mkfs.ext3`, `mkfs.xfs`
 - `mkfs.msdos`, `mkfs.fat`
 - `mkfs -t`
- Typical options:
 - `-b` `blocksize` sets blocksize
 - `-i` `bytes per i-node` sets number of i-nodes
 - `-c` checks disk for bad blocks
- Example:

```
# mkfs.xfs /dev/sda1
...
Writing inode tables: done
Writing superblocks and filesystem accounting info:
done
...
```

Mounting file system

- Using the `mount` command:
 - Supply device name
 - Supply mount point (empty directory)

```
# mount -t xfs /dev/sda1 /mnt/extra
# mount -o ro,remount /dev/sda1 /mnt/extra
```

Optional: Supply file system type

Optional: Supply other options

Optional: Use different superblock

- To show mounted file systems, use `mount` without arguments.

Mounting file systems at system startup

- Add to `/etc/fstab`:

```
/dev/mapper/cl_base-root /      xfs  defaults      2 1
/dev/mapper/cl_base-swap swap swap defaults     1 1
/dev/cdrom    /mnt/cdrom    iso9660  noauto,ro,user 0 0
/dev/fd0      /mnt/floppy   msdos    noauto,user     0 0
/dev/sda6     /mnt/extra    ext3      defaults        0 0
```

- Preferably, using `UUID`:

```
/dev/mapper/cl_base-root /      xfs  defaults      0 0
UUID=2f209047-45fc-43c3-8ab1-8d270ed12b3c /boot ext4  defaults     1 2
/dev/mapper/cl_base-swap swap  swap  defaults     0 0
```

Verify all mounted partitions

- Using commands

```
# mount
```

```
# df [ -k | -m | -h ]
```

```
# lsblk --fs
```

- Using Files

```
# more /etc/mtab
```

```
# more /proc/mounts
```

Unmounting a file system

- Using the `umount` command:
 - disconnect device
 - changing mount options
 - when shutting, temporarily denied access to user

```
# umount /dev/sda1  
# umount /mnt/extra
```

- Unmount all non-os file systems

```
# umount -a
```

- If user still connected

```
# fuser -cu          <- Inform user  
# fuser -ck          <- Kick user out  
# umount -f /mnt
```

The Sequence



- Raw disks : sd | hd | nvme | san luns
- Partitioning : fdisk | parted | gparted
- FS : mkfs | mkfs.xfs | mkfs.ext4
- Mounting : mount /dev/vg/lv /mnt
- Access : cd | touch | mkdir | more | vi | cat ...

Quiz

1. Mount 2nd partition in /dev/sdf to /mnt/data using xfs fs type.
- a) `mount xfs /dev/sd2 /mnt/data`
 - b) `mount -t xfs /dev/sd2 /mnt/data`
 - c) `mount xfs /dev/sdf2 /mnt/data`
 - d) `mount -t xfs /dev/sdf2 /mnt/data`

Quiz

1. Mount 2nd partition in /dev/sdf to /mnt/data using xfs fs type.

- a) mount xfs /dev/sd2 /mnt/data
- b) mount -t xfs /dev/sd2 /mnt/data
- c) mount xfs /dev/sdf2 /mnt/data
- d) mount -t xfs /dev/sdf2 /mnt/data

Quiz

2. You are currently in /mnt/data, and you need to unmount the file system.
- a) `umount -f /mnt/data`
 - b) `kill <pid>; umount /mnt/data`
 - c) `cd; umount /mnt/data`
 - d) `logout then unmount by umount /mntdata cmd`

Quiz

2. You are currently in /mnt/data, and you need to unmount the file system.
- a) `umount -f /mnt/data`
 - b) `kill <pid>; umount /mnt/data`
 - c) `cd; umount /mnt/data`
 - d) `logout` then unmount by `umount /mntdata` cmd

Quiz

3. Why do you need file system?

- a) To provide interface to the hardware
- b) To provide access to storage
- c) To provide interface to the operating system
- d) To provide access to mount point

Quiz

3. Why do you need file system?

- a) To provide interface to the hardware
- b) To provide access to storage
- c) To provide interface to the operating system
- d) To provide access to mount point

Unit summary

Having completed this unit, you should be able to:

- Understand modern internal hard disk type
- Understand Device naming convention
- Perform Disk partitioning
- Understand what and why need File System
- Create and Mount File System
- Unmount File System